

Task4_simple

March 20, 2025

1 Find top 10000 vocab + Logistic Rgression

End of week 10: tasks 0-2

```
[54]: import re
import numpy as np
import pandas as pd
import nltk
from collections import Counter
from nltk.stem import PorterStemmer
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
import requests
import ast

# Download nltk packages if necessary
# nltk.download('punkt')
# nltk.download('stopwords')
```

```
[55]: # ---- 1. Load and prepare data ----

# Mapping of types to "fake" or "reliable"
type_mapping = {
    'unreliable': 'fake',
    'fake': 'fake',
    'conspiracy': 'fake',
    'bias': 'fake',
    'junksci': 'fake',
    'clickbait': 'reliable',
    'reliable': 'reliable',
    'state': 'fake',
    'political': 'reliable',
    'satire': 'fake',
    'hate': 'fake',
```

```

        'rumor': 'fake',
    }

    # Load data and filter relevant columns
    data = pd.read_csv("15,000_rows_preprocessed.csv", usecols=["content", "type"],
        dtype=str)
    #data = pd.read_csv("995,000_rows_preprocessed.csv", usecols=["content",
        "type"], dtype=str)

    # Load BBC-data and add to dataset
    bbc_data = pd.read_csv("BBC_preprocessed.csv", usecols=["content", "type"],
        dtype=str)
    data = pd.concat([data, bbc_data], ignore_index=True)

    # Remove rows with unknown type
    data = data[data['type'] != 'unknown']

    # Map types to labels
    data["label"] = data["type"].map(type_mapping)

    # Remove NaN
    data = data.dropna(subset=["label"])

    # ---- 2. Split dataset to training, validation and test (80/10/10) ----

    train, valid, test = np.split(
        data.sample(frac=1, random_state=42), # Shuffle
        [int(0.8 * len(data)), int(0.9 * len(data))] # Index for split
    )

```

/opt/anaconda3/lib/python3.12/site-packages/numpy/core/fromnumeric.py:59:
FutureWarning: 'DataFrame.swapaxes' is deprecated and will be removed in a
future version. Please use 'DataFrame.transpose' instead.
 return bound(*args, **kwargs)

```

[56]: # Load evaluation data from LIAR-dataset

liar_test_data = pd.read_csv("liar_test_data_preprocessed.csv",
    usecols=['type', 'content'], dtype=str)

type_mapping_liar = {
    'true': 'reliable',
    'false': 'fake',
    'half-true': 'fake',
    'pants-fire': 'fake',
    'barely-true': 'reliable',
    'mostly-true': 'reliable'
}

```

```
}

liar_test_data["label"] = liar_test_data["type"].map(type_mapping_liar)
```

2 Logistic Regression

```
[57]: from sklearn.preprocessing import MaxAbsScaler
import re
import numpy as np
import pandas as pd
import nltk
from collections import Counter
from nltk.stem import PorterStemmer
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report

# ---- 3. Convert text to Bag-of-Words feature matrix ----

content_as_lists = data['content'].apply(ast.literal_eval)

all_words = content_as_lists.explode().tolist()

word_counts = Counter(all_words)

top_10000_words = [word for word, _ in word_counts.most_common(10000)]

vectorizer = CountVectorizer(vocabulary=top_10000_words)
```

```
[58]: # Transform training, valid og test data
X_train = vectorizer.transform(train['content'])
X_valid = vectorizer.transform(valid['content'])
X_test = vectorizer.transform(test['content'])

X_test_liar = vectorizer.transform(liar_test_data['content'])
```

```
[ ]: scaler = MaxAbsScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_valid_scaled = scaler.transform(X_valid)
X_test_scaled = scaler.transform(X_test)

X_test_liar_scaled = (X_test_liar) #scaler.transform
```

```

# Labels
y_train = train['label']
y_valid = valid['label']
y_test = test['label']

y_test_liar = liar_test_data['label']

# ---- 4. Train Logistic Regression model ----

clf = LogisticRegression(max_iter=100000, solver='saga', random_state=42)
clf.fit(X_train_scaled, y_train)

# ---- 5. Evaluate model ----

y_pred_liar = clf.predict(X_test_liar_scaled)

# ---- 6. Evaluate model on LIAR test ----

# y_pred = clf.predict(X_test_liar_scaled)

print(classification_report(y_test_liar, y_pred_liar))

from sklearn.metrics import confusion_matrix

print(confusion_matrix(y_test_liar, y_pred_liar))

# DEBUGGING

print(X_train.shape, X_test_liar.shape)

print(set(y_train), set(y_test_liar))

from collections import Counter
print(Counter(y_train))
print(Counter(y_test_liar))

# Oprindelig test:
print(classification_report(y_test, clf.predict(X_test)))

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| fake | 0.48 | 0.32 | 0.38 | 606 |
| reliable | 0.52 | 0.69 | 0.59 | 661 |
| accuracy | | | 0.51 | 1267 |
| macro avg | 0.50 | 0.50 | 0.49 | 1267 |
| weighted avg | 0.50 | 0.51 | 0.49 | 1267 |

```

[[191 415]
 [207 454]]
(10959, 10000) (1267, 10000)
{'reliable', 'fake'} {'reliable', 'fake'}
Counter({'reliable': 5574, 'fake': 5385})
Counter({'reliable': 661, 'fake': 606})

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| fake | 0.83 | 0.75 | 0.79 | 676 |
| reliable | 0.78 | 0.85 | 0.81 | 694 |
| accuracy | | | 0.80 | 1370 |
| macro avg | 0.80 | 0.80 | 0.80 | 1370 |
| weighted avg | 0.80 | 0.80 | 0.80 | 1370 |